

# STT 465 HW5

*Nate Davis*

*11/30/2017*

## Read in the data

```
data <- read.table("gout.txt")

colnames(data) <- c('sex', 'race', 'age', 'serum_urate', 'gout')

data$gout <- ifelse(data$gout == "Y", 1, 0)
```

## Question 1

### 1.1

```
x <- as.matrix(model.matrix(~sex+age+race, data = data))[, -1]

model <- glm(gout ~ x, data = data, family = 'binomial')
summary(model)
```

```
##
## Call:
## glm(formula = gout ~ x, family = "binomial", data = data)
##
## Deviance Residuals:
##      Min       1Q   Median       3Q      Max
## -0.8294  -0.4256  -0.3537  -0.2713   2.6115
##
## Coefficients:
##              Estimate Std. Error z value Pr(>|z|)
## (Intercept) -8.08409    2.35356  -3.435 0.000593 ***
## xsexM         0.44915    0.38821   1.157 0.247288
## xage          0.09239    0.03644   2.536 0.011227 *
## xraceW       -0.74329    0.41914  -1.773 0.076168 .
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## (Dispersion parameter for binomial family taken to be 1)
##
##      Null deviance: 213.11  on 399  degrees of freedom
## Residual deviance: 203.05  on 396  degrees of freedom
## AIC: 211.05
##
## Number of Fisher Scoring iterations: 5
```

## 1.2

From above we see that the intercept has a very large, very significant effect within the model. The other significant variables at the standard level is age – race is slightly significant but probably negligible.

## 1.3

We have  $Y_i = 0.44915 * 1 + 0.09239 * 65 - 0.74329 * 1 - 8.08409$  which makes our linear predictor equal to about -2.37288. Using our link, we have  $\log(\frac{p_i}{1-p_i}) = -2.37288$ , **making our predicted probability approximately 0.085**

## Question 2

```
# A function to evaluate the log of the posterior density
logP=function(y,X,b,b0,varB){
  Xb=X%*%b
  theta=exp(Xb)/(1+exp(Xb))
  logLik=sum( dbinom(x=y,p=theta,size=1,log=T) )
  logPrior=sum( dnorm(x=b,sd=sqrt(varB),mean=b0,log=T))
  return(logLik+logPrior)
}

logisticRegressionBayes=function(y,X,nIter=12000,V=.02,varB=rep(10000,ncol(X)),b0=rep(0,ncol(X))) {

  ##### Arguments #####
  # y a vector with 0/1 values
  # X incidence matrix fo effects
  # b0,varB, the prior mean and prior variance bj~N(b0[j],varB[j])
  # V the variance of the normal distribution used to generate candidates~N(b[i-1],V)
  # nIter: number of iterations of the sampler
  # Details: generates samples from the posterior distribution of a logistic regression using a
  # Metropolis algorithm
  #####

  # A matrix to store samples
  p=ncol(X)
  B=matrix(nrow=nIter,ncol=p)

  # Centering predictors
  meanX=colMeans(X)
  for(i in 2:p){ X[,i]=(X[,i]-meanX[i]) }

  # A vector to trace acceptance
  accept=rep(NA,nIter)
  accept[1]=TRUE

  # Initialize
  B[1,]=0
  B[1,1]=log(mean(y)/(1-mean(y)))
```

```

    b=B[1,]
    for(i in 2:nIter){

        candidate=rnorm(mean=b,sd=sqrt(V),n=p)

        logP_current=logP(y,X,b0=b0,varB=varB,b=b)
        logP_candidate=logP(y,X,b0=b0,varB=varB,b=candidate)

        r=min(1,exp(logP_candidate-logP_current))
        delta=rbinom(n=1,size=1,p=r)

        accept[i]=delta

        if(delta==1){ b=candidate }
        B[i,]=b

    }

    B[,1]=B[,1]-B[,-1]*%meanX[-1] # absorbing means on the intercept

    return(list(B=B,accept=accept))
}

x_new = model.matrix(~sex+age+race,data=data)

b_model <- logisticRegressionBayes(data$gout, x_new, 55000)

result <- cbind(model$coefficients,colMeans(b_model$B[-(1:5000),]))

params <- b_model$B[-(1:5000),]
colnames(params) <- c("Intercept", "Sex", "Age", "Race")

```

## 2.1

```

library("coda")

## Warning: package 'coda' was built under R version 3.3.2
params <- data.frame(params)

mcmc_obj <- mcmc(params)

obj_sum <- summary(mcmc_obj)

obj_sum

##
## Iterations = 1:50000
## Thinning interval = 1
## Number of chains = 1
## Sample size per chain = 50000
##
## 1. Empirical mean and standard deviation for each variable,

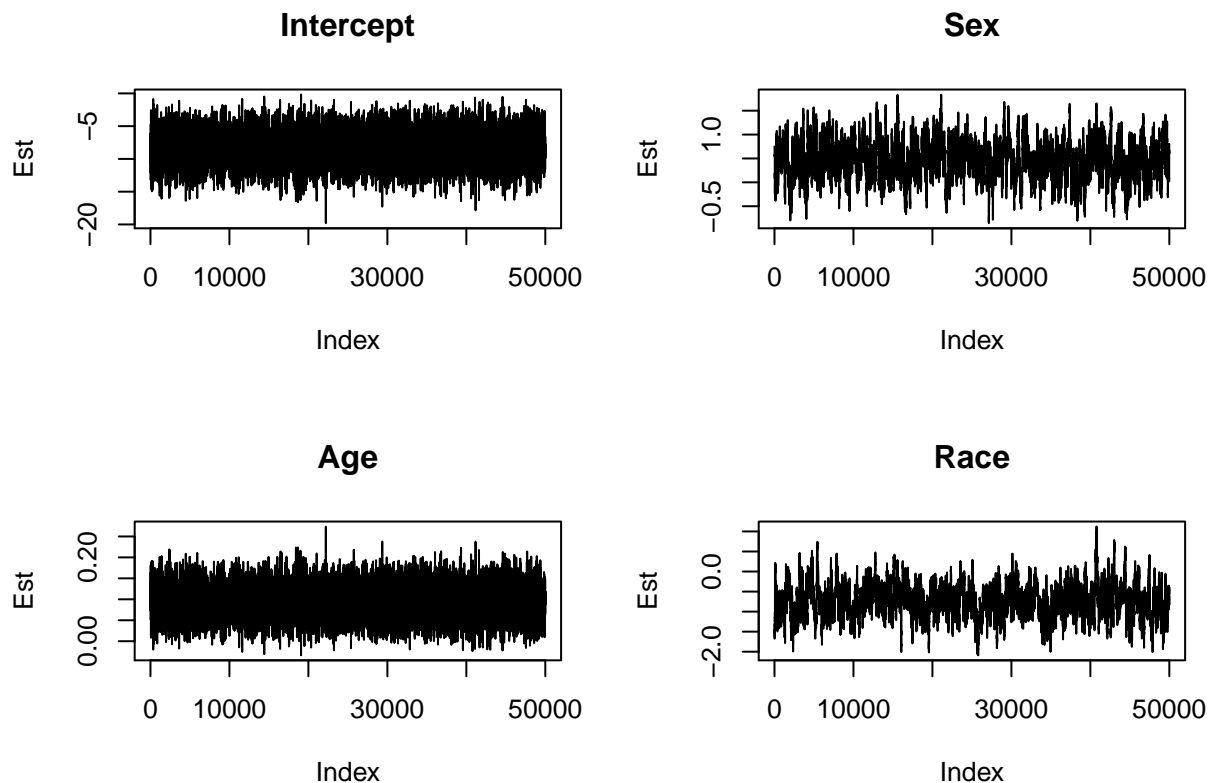
```

```
## plus standard error of the mean:
##
##           Mean      SD Naive SE Time-series SE
## Intercept -8.33091 2.37854 0.0106372      0.037958
## Sex        0.46491 0.40030 0.0017902      0.021442
## Age        0.09514 0.03694 0.0001652      0.000565
## Race       -0.73726 0.42329 0.0018930      0.024159
##
## 2. Quantiles for each variable:
##
##           2.5%      25%      50%      75%      97.5%
## Intercept -13.11548 -9.9164 -8.30724 -6.7052 -3.8026
## Sex        -0.35085  0.1954  0.46686  0.7360  1.2452
## Age         0.02378  0.0702  0.09497  0.1199  0.1686
## Race       -1.53226 -1.0322 -0.74542 -0.4608  0.1242
```

## 2.2

```
par(mfrow = c(2, 2))

plot(params$Intercept, xlab = "Index", ylab = "Est", main = "Intercept", type = "l")
plot(params$Sex, xlab = "Index", ylab = "Est", main = "Sex", type = "l")
plot(params$Age, xlab = "Index", ylab = "Est", main = "Age", type = "l")
plot(params$Race, xlab = "Index", ylab = "Est", main = "Race", type = "l")
```



```
size_se <- cbind(effectiveSize(mcmc_obj), obj_sum$statistics[,3])
colnames(size_se) <- c("Effective Size", "MC SE")
```

```
size_se
```

##	Effective Size	MC SE
## Intercept	3926.5313	0.0106371562
## Sex	348.5334	0.0017901858
## Age	4274.2771	0.0001652075
## Race	306.9899	0.0018929985

## 2.3

```
pts <- cbind(rep(c(0, 1), each = 4), rep(c(55, 65), 4), rep(c(1, 0), 2, each = 2))

par(mfrow = c(3, 3))

for(i in 1:length(pts[,1])){

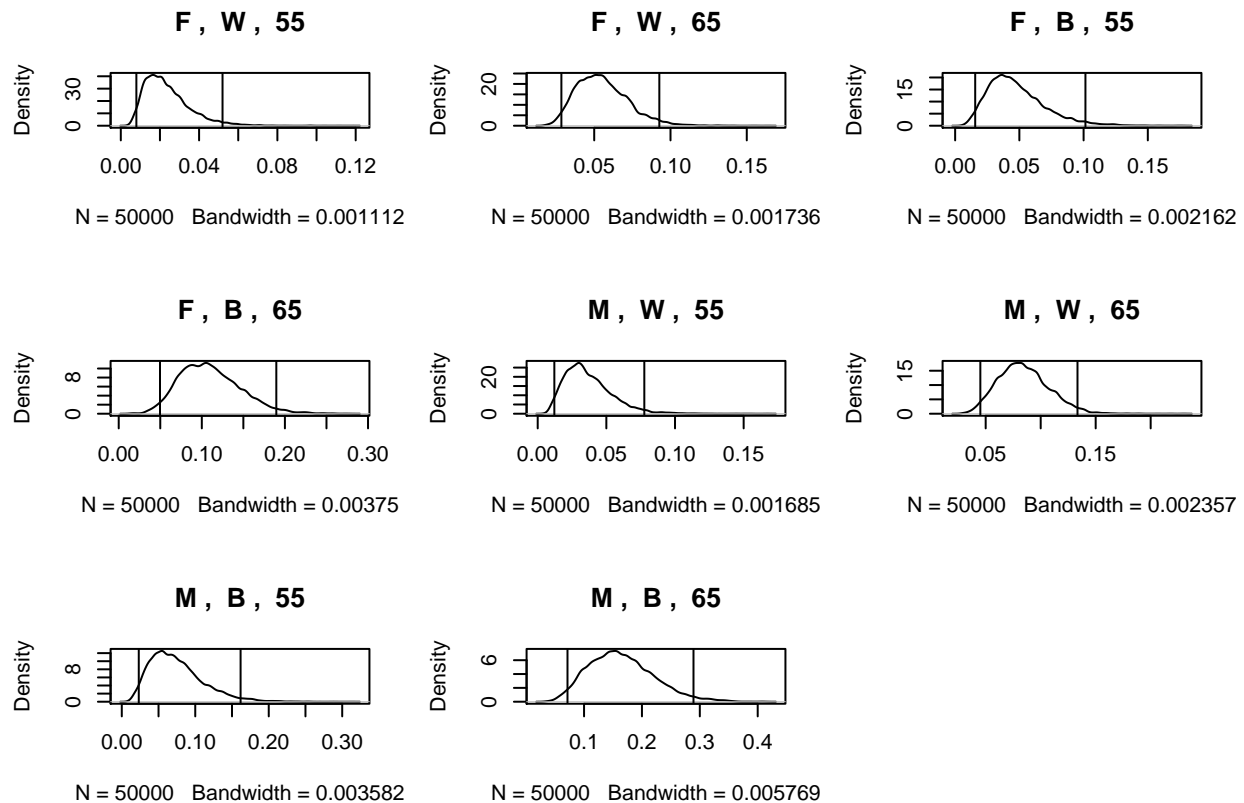
  resp <- params[, 1] + params[, 2]*pts[i, 1] + params[, 3]*pts[i, 2]+ params[, 4]*pts[i, 3]

  resp <- exp(resp) / (exp(resp) + 1)

  sex <- ifelse(pts[i, 1] == 0, "F", "M")
  race <- ifelse(pts[i, 3] == 0, "B", "W")

  title <- paste(sex, ", ", race, ", ", as.character(pts[i, 2]))

  plot(density(resp), main = title)
  abline(v = quantile(resp, probs = c(0.025, 0.975)))
}
```



## Question 3

### 3.1

```
v_values <- c(.4, .2, .05, .001)

res <- cbind(rep(rep(NA, 4)), NA, NA)
row.names(res) <- v_values
colnames(res) <- c("Acc Rate", "Lag Corr", "Effective Size")

for(i in 1:length(v_values)){

  b_model <- logisticRegressionBayes(data$gout, model.matrix(~sex+age+race,data=data),
                                     55000, V = v_values[i])

  params <- b_model$B[-(1:5000),]
  colnames(params) <- c("Intercept", "Sex", "Age", "Race")

  res[i, 1] <- sum(b_model$accept)/length(b_model$accept)

  obj <- mcmc(params)

  res[i, 2] <- autocorr(obj[,3], lags = 100)

  res[i, 3] <- effectiveSize(obj)[3]
```

```
}
```

```
res
```

##		Acc Rate	Lag Corr	Effective Size
##	0.4	0.01843636	0.11222517	665.3349
##	0.2	0.03587273	0.06270784	1112.9198
##	0.05	0.12898182	-0.01003891	3653.8466
##	0.001	0.70643636	0.09698724	1781.4365

## 3.2

The best value from above should be around 0.05 as it not only has the consistently smallest lag correlation value, but also the greatest effective sample size. We get more information from our algorithm when its ran with  $V = 0.05$ .